

## Nvidia Titan V - Real time object detection

**Notebook:** 1.Getting Started

**Created:** 4/9/2018 11:01 AM

**Updated:** 20/2/2019 10:11 AM

**Author:** pengpeng

**Tags:** CV, DL, GPU, TF

---

This article is briefing about my experience on setting real time object detection using Nvidia Titan V.

Here highlighted the working combination of software version and step.

### Hardware required:

1. Workstation - Dell PowerEdge T320
2. Nvidia Graphic Card - Titan V

### Software required :

1. OS - Ubuntu 18.04 Bionic Beaver
2. Linux kernel header - 4.16.0-041600-generic
3. Nvidia Driver - 396.44
4. CUDA release version - 9.2
5. Cudnn - v7.2.1
6. TensorFlow 1.8.0
7. OpenCV 3.4.2

### Step 1 - Setup Linux kernel :

Install ubuntu 18.04 on the workstation. By default, ubuntu 18.04, we get "4.15.0-23-generic". Upgrade kernel to 4.16 version as command :

#### # Download linux kernel files :

```
$ mkdir -p ~/kernel-4_16 && cd ~/kernel-4_16
$ wget http://kernel.ubuntu.com/~kernel-ppa/mainline/v4.16/linux-headers-4.16.0-041600_4.16.0-041600.201804012230_all.deb
$ wget http://kernel.ubuntu.com/~kernel-ppa/mainline/v4.16/linux-headers-4.16.0-041600-generic_4.16.0-041600.201804012230_amd64.deb
$ wget http://kernel.ubuntu.com/~kernel-ppa/mainline/v4.16/linux-image-4.16.0-041600-generic_4.16.0-041600.201804012230_amd64.deb
files
```

#### # Install linux-headers :

```
$ sudo dpkg -i *.deb
```

#### # After installation is complete, reboot the system and Verify :

```
$ uname -sr
```

#### # Output :

```
kpo@kpo-T320: ~  
File Edit View Search Terminal Help  
kpo@kpo-T320:~$ uname -r  
4.16.0-041600-generic  
kpo@kpo-T320:~$
```

## Step 2 - Install Nvidia - Cuda toolkit

# Remove previous cuda installation(if you installed cuda before):

```
$ sudo apt-get purge nvidia*  
$ sudo apt-get autoremove  
$ sudo apt-get auto clean  
$ sudo rm -rf /usr/local/cuda*
```

# Install cuda:

```
$ sudo apt-key adv --fetch-keys  
http://developer.download.nvidia.com/compute/cuda/repos/ubuntu1710/x86_64/7fa2af80.pub  
$ echo "deb  
https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1710/x86_64 /" | sudo tee  
/etc/apt/sources.list.d/cuda.list  
$ sudo apt-get update  
$ sudo apt-get -o Dpkg::Options::="--force-overwrite" install cuda-9-2 cuda-drivers  
cuda=9.0.176-1
```

# After installation is complete, reboot the system to reload driver.

# Setup linux environment for CUDA compilation :

```
$ sudo apt-get install libcupti-dev  
/usr/local/cuda-9.2/bin  
echo 'export PATH=/usr/local/cuda-9.2/bin:$PATH' >> ~/.bashrc
```

```
$ echo 'export LD_LIBRARY_PATH=/usr/local/cuda-  
9.2/lib64${LD_LIBRARY_PATH:+:${LD_LIBRARY_PATH}}' >> ~/.bashrc  
$ source ~/.bashrc  
$ sudo ldconfig
```

# Verify Nvidia-driver and Cuda version:

```
$ nvidia-smi  
$ nvcc -V
```

# Output :

```
kpo@kpo-T320: ~  
File Edit View Search Terminal Help  
kpo@kpo-T320:~$ nvidia-smi  
Mon Sep  3 17:57:55 2018  
+-----+  
| NVIDIA-SMI 396.44                  Driver Version: 396.44 |  
+-----+  
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |  
| Fan  Temp  Perf    Pwr:Usage/Cap|     Memory-Usage | GPU-Util  Compute M. |  
+-----+  
|    0  TITAN V         Off         | 00000000:0A:00.0 Off |                    |  
| 30%   45C   P0      37W / 250W |  0MiB / 12066MiB |      0%    Default |  
+-----+  
+-----+  
| Processes:                                                       GPU Memory |  
|  GPU       PID    Type    Process name                       Usage    |  
+-----+  
| No running processes found                                     |  
+-----+  
kpo@kpo-T320:~$ nvcc -V  
nvcc: NVIDIA (R) Cuda compiler driver  
Copyright (c) 2005-2018 NVIDIA Corporation  
Built on Tue_Jun_12_23:07:04_CDT_2018  
Cuda compilation tools, release 9.2, V9.2.148  
kpo@kpo-T320:~$
```

# Install GPU-accelerated library cuDNN v7.2.1:

#Download cudnn-9.2-linux-x64-v7.2.1.38.tgz from <https://developer.nvidia.com/cudnn>

\$ cd ~/Downloads #assume file is saved in folder "~/Downloads"

\$ tar -xvf cudnn-9.2-linux-x64-v7.2.1.38.tgz

\$ sudo cp -R cuda/include/\* /usr/local/cuda-9.2/include

\$ sudo cp -R cuda/lib64/\* /usr/local/cuda-9.2/lib64

#Install Dependencies (**libcupti**):


\$ sudo apt-get install libcupti-dev

\$ echo 'export LD\_LIBRARY\_PATH=/usr/local/cuda/extras/CUPTI/lib64:\$LD\_LIBRARY\_PATH' >>

~/.bashrc

### Step 3 : Instal opencv 3.4.2

# Run the script below to install opencv 3.4.2

 [install-opencv.sh](#)  
1/4/2019 5:23 PM, 2.4 KB

Step 4 : Install TensorFlow v1.8 by source :

# Install dependencies :

\$ sudo apt-get install python-numpy python-dev python-pip python-wheel

\$ sudo apt-get install python3-numpy python3-dev python3-pip python3-wheel

# Install Bazel - to build tensorflow source file:

\$ cd ~/

```
$ wget https://github.com/bazelbuild/bazel/releases/download/0.14.0/bazel-0.14.0-installer-linux-x86\_64.sh
$ chmod +x bazel-0.14.0-installer-linux-x86_64.sh
$ ./bazel-0.14.0-installer-linux-x86_64.sh --user
$ echo 'export PATH="$PATH:$HOME/bin"' >> ~/.bashrc
# Reload environment variables
$ source ~/.bashrc
$ sudo ldconfig
# Download tensorflow 1.8.0 and configure
$ cd ~/
$ git clone https://github.com/tensorflow/tensorflow.git
$ cd tensorflow
$ git pull
$ git checkout r1.8
$ ./configure
# tensorflow configuration as :
```

```

kpo@kpo-T320:~$ cd ~/tensorflow/
kpo@kpo-T320:~/tensorflow$ ./configure
You have bazel 0.14.0 installed.
Please specify the location of python. [Default is /usr/bin/python]: /usr/bin/python3

Found possible Python library paths:
  /usr/lib/python3/dist-packages
  /usr/local/lib/python3.6/dist-packages
Please input the desired Python library path to use. Default is [/usr/lib/python3/dist-pa
ckages]

Do you wish to build TensorFlow with jemalloc as malloc support? [Y/n]: Y
jemalloc as malloc support will be enabled for TensorFlow.

Do you wish to build TensorFlow with Google Cloud Platform support? [Y/n]: Y
Google Cloud Platform support will be enabled for TensorFlow.

Do you wish to build TensorFlow with Hadoop File System support? [Y/n]: Y
Hadoop File System support will be enabled for TensorFlow.

Do you wish to build TensorFlow with Amazon S3 File System support? [Y/n]: Y
Amazon S3 File System support will be enabled for TensorFlow.

Do you wish to build TensorFlow with Apache Kafka Platform support? [Y/n]: n
No Apache Kafka Platform support will be enabled for TensorFlow.

Do you wish to build TensorFlow with XLA JIT support? [y/N]: n
No XLA JIT support will be enabled for TensorFlow.

Do you wish to build TensorFlow with GDR support? [y/N]: n
No GDR support will be enabled for TensorFlow.

Do you wish to build TensorFlow with VERBS support? [y/N]: n
No VERBS support will be enabled for TensorFlow.

Do you wish to build TensorFlow with OpenCL SYCL support? [y/N]: n
No OpenCL SYCL support will be enabled for TensorFlow.

Do you wish to build TensorFlow with CUDA support? [y/N]: y
CUDA support will be enabled for TensorFlow.

Please specify the CUDA SDK version you want to use, e.g. 7.0. [Leave empty to default to
CUDA 9.0]: 9.2

Please specify the location where CUDA 9.2 toolkit is installed. Refer to README.md for mo
re details. [Default is /usr/local/cuda]:

Please specify the CUDA SDK version you want to use, e.g. 7.0. [Leave empty to default to
CUDA 9.0]: 9.2

Please specify the location where CUDA 9.2 toolkit is installed. Refer to README.md for mo
re details. [Default is /usr/local/cuda]:

Please specify the cuDNN version you want to use. [Leave empty to default to cuDNN 7.0]: 7
.2.1

Please specify the location where cuDNN 7 library is installed. Refer to README.md for mor
e details. [Default is /usr/local/cuda]:

Do you wish to build TensorFlow with TensorRT support? [y/N]: n
No TensorRT support will be enabled for TensorFlow.

Please specify the NCCL version you want to use. [Leave empty to default to NCCL 1.3]:

Please specify a list of comma-separated Cuda compute capabilities you want to build with.
You can find the compute capability of your device at: https://developer.nvidia.com/cuda-gpus.
Please note that each additional compute capability significantly increases your build tim
e and binary size. [Default is: 7.0]

Do you want to use clang as CUDA compiler? [y/N]: n
nvcc will be used as CUDA compiler.

Please specify which gcc should be used by nvcc as the host compiler. [Default is /usr/bin

```

```
/gcc]:

Do you wish to build TensorFlow with MPI support? [y/N]: n
No MPI support will be enabled for TensorFlow.

Please specify optimization flags to use during compilation when bazel option "--config=opt" is specified [Default is -march=native]:

Would you like to interactively configure ./WORKSPACE for Android builds? [y/N]: n
```

# Configuration finished

# Build tensorflow using bazel \*This process takes a fairly long time.

```
$ bazel build --config=opt //tensorflow/tools/pip_package:build_pip_package
```

# To build whl file issue following command:

```
$ bazel-bin/tensorflow/tools/pip_package/build_pip_package tensorflow_pkg
```

# Install tensorflow with pip:

```
$ cd tensorflow_pkg
```

```
$ pip3 install tensorflow*.whl
```

# Verify tensorflow installation

- Run in terminal

**# output as :**

```
python3
import tensorflow as tf
hello = tf.constant('Hello, TensorFlow!')
sess = tf.Session()
print(sess.run(hello))
```

```
kpo@kpo-T320:~/tensorflow/tensorflow_pkg$ python3
Python 3.6.5 (default, Apr 1 2018, 05:46:30)
[GCC 7.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import tensorflow as tf
>>> hello = tf.constant('Hello, TensorFlow!')
>>> sess = tf.Session()
2018-09-03 19:08:40.953059: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1356] Found device 0 with properties:
name: TITAN V major: 7 minor: 0 memoryClockRate(GHz): 1.455
pciBusID: 0000:0a:00:0
totalMemory: 11.78GiB freeMemory: 11.37GiB
2018-09-03 19:08:40.953120: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1435] Adding visible gpu devices: 0
2018-09-03 19:08:41.422234: I tensorflow/core/common_runtime/gpu/gpu_device.cc:923] Device interconnect StreamExecutor
with strength 1 edge matrix:
2018-09-03 19:08:41.422280: I tensorflow/core/common_runtime/gpu/gpu_device.cc:929] 0
2018-09-03 19:08:41.422301: I tensorflow/core/common_runtime/gpu/gpu_device.cc:942] 0: N
2018-09-03 19:08:41.422732: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1053] Created TensorFlow device (/job:lo
calhost/replica:0/task:0/device:GPU:0 with 10997 MB Memory) -> physical GPU (device: 0, name: TITAN V, pci bus id: 0000
:0a:00:0, compute capability: 7.0)
>>> sess = tf.Session()
2018-09-03 19:08:47.994200: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1435] Adding visible gpu devices: 0
2018-09-03 19:08:47.994290: I tensorflow/core/common_runtime/gpu/gpu_device.cc:923] Device interconnect StreamExecutor
with strength 1 edge matrix:
2018-09-03 19:08:47.994311: I tensorflow/core/common_runtime/gpu/gpu_device.cc:929] 0
2018-09-03 19:08:47.994329: I tensorflow/core/common_runtime/gpu/gpu_device.cc:942] 0: N
2018-09-03 19:08:47.994639: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1053] Created TensorFlow device (/job:lo
calhost/replica:0/task:0/device:GPU:0 with 10997 MB Memory) -> physical GPU (device: 0, name: TITAN V, pci bus id: 0000
:0a:00:0, compute capability: 7.0)
>>>
```

## Step 5 : Real time Object Detection - tensorflow - CoCo Model

```
$ cd ~/tensorflow
```

```
$ git clone https://github.com/tensorflow/models.git
```

```
$ sudo apt-get install protobuf-compiler python-pil python-lxml python-tk
```

```
$ pip3 install --user Cython
```

```
$ pip3 install --user contextlib2
```

```
$ pip3 install --user pillow
```


```
$ pip3 install --user lxml
```

```
$ pip3 install --user jupyter
```

```
$ pip3 install --user matplotlib
$ cd ~/tensorflow/models/research
$ protoc object_detection/protos/*.proto --python_out=.
$ export PYTHONPATH=$PYTHONPATH:`pwd`:`pwd`/slim
```

```
$ cd object_detection
```

# Place 'object\_detection\_webcam.py' in ~/tensorflow/models/research/object\_detection directory

 object\_detection\_webcam.py  
1/4/2019 5:23 PM, 4.2 KB

```
$ python3 object_detection_webcam.py
# output as :
```

